

Feature Engineering 101



Dan Ofer

<https://ddofer.github.io>

Overview

1. What is feature engineering (FE)? Why is it important for machine learning / statistics?
2. Feature engineering techniques for:
 - Categorical features
 - Text (Natural language)
 - Time Series
 - Geospatial

Why listen to me?

- **Dan Ofer** - Senior Data Scientist 4.5Y, now at Nutrino/Medtronic
- [Sparkbeyond](#): AI with Fortune 500 & charities, including healthcare (Clalit), insurance, churn, chemicals etc'
- Top 0.8% on Kaggle (kaggle.com/danofer)
- 1st Place in [WiDS 2020 challenge](#)
- MsC: Neuroscience & Bioinformatics, [thesis on protein feature engineering](#) (HUJI)
- Probably took your picture at a convention/Midburn!




Feature Engineering

"Applied machine learning" is basically feature engineering

- Andrew Ng


Feature Engineering: What is it good for?



Feature Engineering: What is it good for?

Coordinate transformation


$$r = \sqrt{x^2 + y^2} \quad \theta = \arctan \frac{y}{x}$$



Tangled




Feature engineering



Transparent

**Feature engineering =
transforming raw data into
features that better
represent the underlying
problem, resulting in
improved predictive model
accuracy on unseen data**

Feature Engineering



- Make *good* features
 - Highly predictive (of the target variable), succinct, interpretable, robust
- “Secret sauce of machine learning”
 - ~3d most important part of an machine learning project (after problem definition & data cleaning)
- Related, but different:
 - Feature extraction (raw features)
 - Feature selection
- Huge topic. More art than science
 - Extremely manual, handcrafted “dark magic” (For now)
- [SparkBeyond](#) (SOTA Automated Feature Engineering)
- Further Reading: [Appendix + A few useful things to Know about machine Learning](#)

Feature Extraction Vs Engineering

Given a Text, predict emotional sentiment

- “Star Wars = Greatest. movie. Ever!! Awesome! :D “ : Positive
- “Coronavirus quarantine/furlough makes me sad+bored” : Negative

- Feature extraction:

- Characters or bytes in text

- (Simple) Feature engineering:

- Lowercase text, tokenize (“split”) by whitespace & punctuation, count word frequencies etc’
 - What’s a word? What about Chinese? What about “New-York”?

Categorical Features

Categorical Features

- Categorical or ordinal features. “Moderate” cardinality (lower than free text, most elements expected to appear multiple times)
- **Cardinality** = How many unique values
- **Ordinal** features - have explicit “order”: $3 > 2 > 1$. Old > Young...
 - Gender - [M/F/Rather not say].
 - Education level {Ordinal}. [Primary school/HS/college/graduate/PhD]
 - Ad category. [Sports/music/politics/cars/...]
 - Product ID. [Fuji XT3 / iPhone 10 / Galaxy S20/...]
- Some ML libraries handle “automatically”! E.g. Catboost, LGBM

ONE-HOT ENCODING

Feature		Apple	Pear
Apple	ONE HOT ENCODING	1	0
Pear		0	1
Apple		1	0
Pear		0	1
Apple		1	0

One-hot encoding allows us to turn nominal categorical data into features with numerical values, while not mathematically imply any ordinal relationship between the classes.

ChrisAlbon

One Hot Encoding (OHE)

- Naive - **One Hot Encoding** - create a column for each possible value
- Can result in a lot of columns, memory use. Can work best with simple models and huge amounts of data. Not ideal for tree models, when high cardinality.
- More approaches - [feature hashing](#), etc'

One hot encoding methods:

```
df = pd.get_dummies(df, drop_first=True)
```

```
>>> from sklearn.preprocessing import OneHotEncoder
```

```
>>> X = [['Male', 1], ['Female', 3], ['Female', 2]]
```

```
>>> enc = OneHotEncoder(handle_unknown='ignore').fit(X)
```

```
>>> enc.categories_
```

```
[array(['Female', 'Male'], dtype=object), array([1, 2, 3], dtype=object)]
```

Ordinal encoding

- Replace string with (integer) number
- Use scikit-learn's [OrdinalEncoder](#)
- Can also use Pandas Dataframe's Categorical type - supports ordinals/ordering
- Works well for tree models that can “cut” arbitrarily. Unsuitable for linear models

```
>>> from sklearn.preprocessing import OrdinalEncoder, LabelEncoder
```

```
>>> le = LabelEncoder()
```

```
>>> le.fit(["paris", "paris", "tokyo", "amsterdam"])
```

```
>>> le.transform(["tokyo", "tokyo", "paris"])
```

```
array([2, 2, 1]...)
```

Ordinal/Label vs One-Hot Encoding

Label Encoding

Food Name	Categorical #
Apple	1
Chicken	2
Broccoli	3



One Hot Encoding

Apple	Chicken	Broccoli
1	0	0
0	1	0
0	0	1

Big Idea - reduce dimensionality:

- **Reduce # unique variables to learn**
- **Handle rare variables (e.g. “singletons”)**

Transformations: Count/Frequency Encoding

- Replace variables with their count in the data (“[Count encoding](#)”)
 - Combo! Replace variables that appear less than K times with their count.
- Efficient for reducing cardinality - e.g. if 70% of our values appear less than 2 times in the data.
- Relatively robust
- Adds some new information! (Popular/rare values)

```
data.apply(lambda x: x.map(x.value_counts()))
```

```
data.where(data.apply(lambda x: x.map(x.value_counts()))>=2, "other")
```

Transformations: Target Encoding

- Replace variables with target frequency - target/label encoding
- Some smoothing must be used, as otherwise will overfit badly!
- Many [variants](#) to technique: [Weight Of Evidence](#), Bayesian target encoding, use of [nested cross validation](#), etc'
 - Catboost does this internally
- Gotcha: just adding smoothing won't help for variables that appear just 1 time - they'll still be overfit, no matter the global prior!

```
df["category_average_target"] = (0.5 +  
df.groupby(["category"])[["target"]].mean())/2
```

Interactions/Feature Crossings

- Statistical features over groups, or of a feature vs the group (e.g. an item's price vs mean price of all items in store-department).

```
df['Mean_Category_Price'] =  
df.groupby(["Item_category"])[ 'Price' ].transform("mean")  
  
df['relative_price'] = df['Price'].div(df['Mean_Category_Price'])
```

- Conjoined features:
 - Useful when we need a “unique ID” to distinguish an entity
 - E.g.: predict sales per item, per store:

```
df["joint_id"] = df["product_id"] + df["store_id"]
```

Interactions/Feature Crossings

Combinations (“crosses”) of multiple features - “A & B !C”


- e.g. “Age <20 & Education < college & profession == Doctor” -> Fraud
 - sklearn's [PolynomialFeatures](#)
 - Can combine with bucketing and feature hashing -
https://www.tensorflow.org/tutorials/structured_data/feature_columns#crossed_feature_columns
- Can easily overfit and create huge amount of noise features! Feature selection may not be enough
- I recommend restrict functions and columns used for generating features
- Consider your ml model's expressiveness, and if interaction expressed a non-linear relationship - e.g. BMI (Height/Weight²)
 - This probably works well with huge amounts of data and regularized linear models

“Rounding down”

- Reduce cardinality (# unique variables) by aggregating to a higher “level”.
 - Reduces information (Bad 😞)
 - Reduces cardinality (good) -> reduces overfitting (good 😊). Domain specific...
- How to do it? - Take Prefix (truncate), round down granularity...
 - IP Addresses - take first few “blocks”.
“192.168.42.1” -> “192.168”
Map IPs to Country/state/city/zipcode..
 - Software version:
“Windows 10.1.17.2358” -> “Windows 10.1.17”
 - Emails, websites - extract the domain.
www.mashable.com/news/dan-ofer-wins-nobel-for-extreme-cleverness -> “mashable”
 - Hierarchical codes - zipcodes, medical codes (ICD9..) - truncate first X digits.
Zipcode “90210” - Beverly hills neighbourhood. “902” -~ County level (larger area)
 - Age: 24.16 -> 24
 - Works best with semantic understanding

Entity Embedding

- Embedding using linear algebra/ML model. Word2Vec style
- Can also use SVD, matrix decomposition, graph methods, etc’.
- Used in Kaggle Rossman competition - “entity embedding”
 - Entity Embeddings of Categorical Variables (2016)
- Won’t help with singletons/rare variables, unless you have a lot of unsupervised data to use. Can underperform simple one hot encoding - depends on ML model, variable cardinality, data size.



Predict emotional sentiment of a Text

- “Star Wars = Greatest. movie. Ever!! Awesome! :D” : **Positive**
- “Coronavirus quarantine/furlough makes me sad+bored” : **Negative**

www.kaggle.com/danofer/reddit-comments-scores-nlp

<https://github.com/ddofer/talk/blob/master/NLP%20101%20-%20ML%20Seminar%202017.pdf>

Text Features

Text/NLP preprocessing

Huge subject in itself - transforming the input text before getting features.

- Lowercase text (or don't!). (“DAN Is grEAT” -> “dan is great”)
- Normalize words, contractions, phrases, acronyms. E.g.:
 - “it’s” -> “it is”
 - “{England, GB, Blighty}” -> “England”.
 - Spelling correction
- Drop, keep or substitute placeholders for entities, phone numbers, emails..
“Call 052-9021042” -> “call PHONE_NUM”
- Tokenizers (custom delimiters. E.g. Twitter tokenizer..).
- Stem or Lemmatizers: “Cats” -> “Cat” ; “Octopii”, “Octopuses” -> “Octopus”
- Stop word removal - (“for and the or I”)

Excellent tools: [Spacy](#), [Textacy](#), [Scikit-learn](#), [Gensim](#), [NLTK](#)

https://chartbeat-labs.github.io/textacy/build/html/api_reference/text_processing.html

Bag of Words = count words in text

Document 1

The quick brown
fox jumped over
the lazy dog's
back.

Document 2

Now is the time
for all good men
to come to the
aid of their party.

Term	Document 1	Document 2
aid	0	1
all	0	1
back	1	0
brown	1	0
come	0	1
dog	1	0
fox	1	0
good	0	1
jump	1	0
lazy	1	0
men	0	1
now	0	1
over	1	0
party	0	1
quick	1	0
their	0	1
time	0	1

Stopword List

for
is
of
the
to

Improving on Bag of Words: TF/TF-IDF

- Bag of Words = count each words in text [0/1]
- TF - Term frequency (How many times a word appeared in the text)
- TF-IDF (IDF = Inverse document frequency: give less weight to very frequent words)

https://scikit-learn.org/stable/tutorial/text_analytics/working_with_text_data.html

https://chrisalbon.com/machine_learning/preprocessing_text/bag_of_words/

```
from sklearn.feature_extraction.text import CountVectorizer  
from sklearn.feature_extraction.text import TfidfTransformer
```

TF-IDF

TF-IDF is a measure of originality of a word by comparing the number of times a word appears in a doc with the number of docs the word appears in.

$$\text{TF-IDF} = \text{TF}(t, d) \times \text{IDF}(t)$$

Term frequency

Number of times term t appears in a doc, d

Inverse document frequency

$$\log \frac{1 + \overset{\text{\# of documents}}{n}}{1 + \underset{\text{Document frequency of the term } t}{df(d, t)}}$$

Improving on Bag of Words: TF/TF-IDF

- Bag of Words = count each words in text [0/1]
- TF - Term frequency (How many times a word appeared in the text)
- TF-IDF (IDF = Inverse document frequency: downweighs common words)

My advice: Try this first (with defaults and lowercasing).

$$w_{x,y} = \text{tf}_{x,y} \times \log \left(\frac{N}{\text{df}_x} \right)$$

TF-IDF

Term x within document y

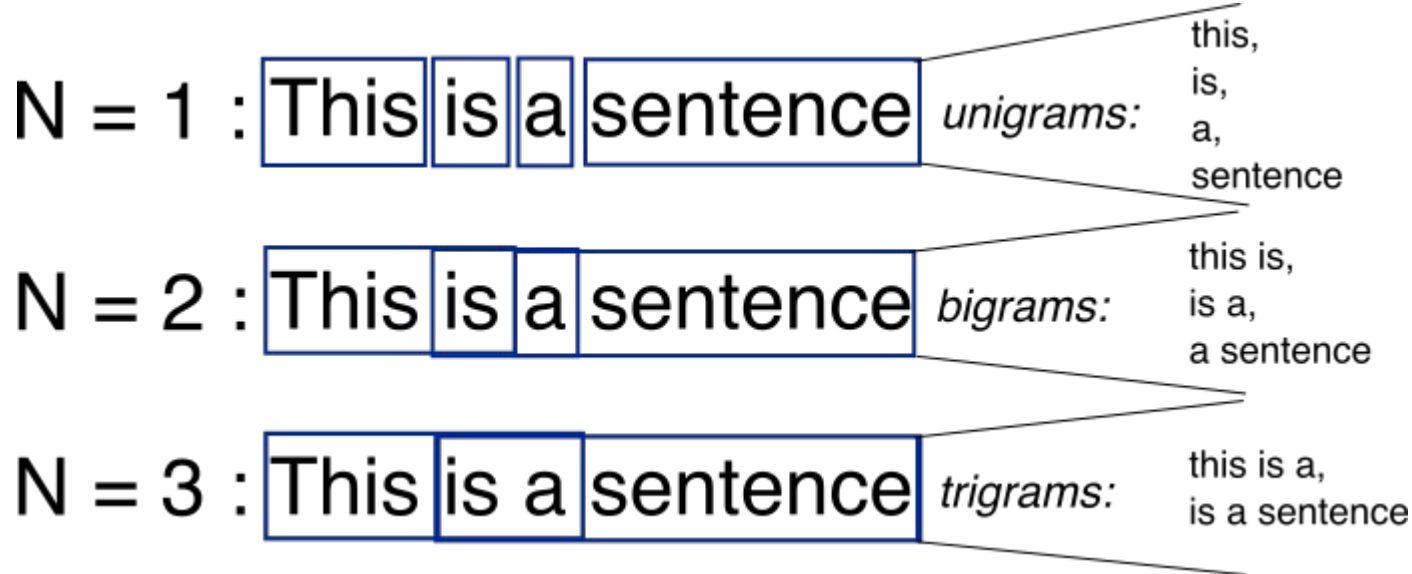
$\text{tf}_{x,y}$ = frequency of x in y

df_x = number of documents containing x

N = total number of documents


N-grams

- N-grams - N words together
 - Unigrams (1-gram/"default") - "New York" -> {"New":1, "York":1}
 - 2-Grams: "New York" -> {"New York":1}



Character level N-grams (shingles)

- Use letters/characters as “tokens”/basic unit, instead of words
- **Character level n-grams (shingles)** can be VERY useful for some domains!
 - Gives “free” count of special characters, punctuation etc’ (#,@,!, \$...)
 - E.g. 3-gram characters shingles: “Danny” -> {“dan”, “ann”, “nny” }



```
from sklearn.feature_extraction.text import CountVectorizer
vectorizer = CountVectorizer(analyzer="char", ngram_range=(4,4))
```

https://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.CountVectorizer.html

Basic Text Featurization - tips

- EDA: frequency of majority/minority classes, text length, most frequent words, empty sentences, duplicates, html crud etc'
- Using N-grams - limit max N-size, max vocabulary, consider iterative text cleaning (e.g. conjoin entities ("New_York") in advance, remove stop-words)
- Don't "blindly" drop stop words in advance
- TF-IDF/BoW - try counts or frequency. Min count 3-5. Max_df ~ 0.98 %
- CountVectorizer also useful to featurize categorical "like" features (e.g. list of entities)
- Note - sklearn vectorizers use sparse matrices (Pandas supports them)
- Feature selection (e.g. Chi2, mutual information, max vocab size)
- SVD on term/BoW matrix

Text features

- Phrases/coallocations (Gensim; “New_York”) - expand n-gram space cheaply
- Semantic attributes - reading level (Textacy), named entities (Company, date, location etc’), Parts of speech (e.g. nouns, adjectives), % Camel Cased, % ALL CAPITALS, words in a predefined list (profanity)
- Language model - score sentence likelihood with [KenLM](#) (fast), Bert/GPT NN etc’
- Named entity extraction (NER - Spacy) - extract entities, and features about them:
E.g. “contains phone number -> Number is in New York area”.

My example features notebook with code:


<https://www.kaggle.com/danofer/reddit-comments-scores-nlp#Text-features-engineering>

Word2Vec Text Features

- Word2Vec/FastText/Glove/Doc2Vec embeddings - get mean/max/sum over entire text, use vectors as new features
 - Train from scratch or use [pre-trained embeddings](#).
 - Positional NN model embeddings like BERT - won't necessarily be better!
 - Multiple by word level TF-IDF score to improve - “[A simple but tough to beat baseline](#)”
 - Consider domain specific pretrained embeddings, +- fine-tuning.
 - Fast to train (unlike BERT)
 - Recommended library - many, but start with [Gensim](#)

Advanced NLP augmentation - External data

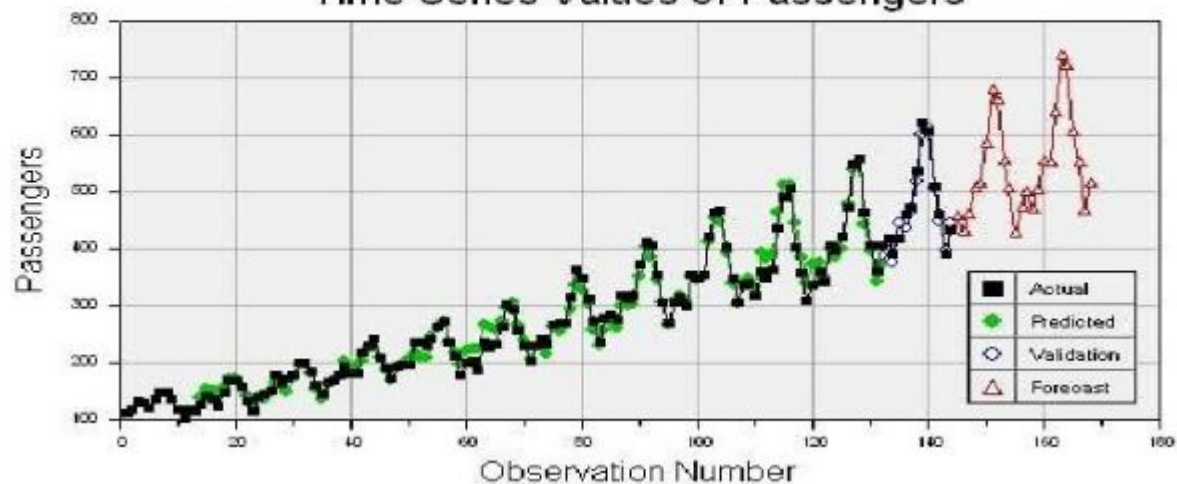
- Simple: Emotion/Sentiment Lexicons.
 - {"Great":+3, good:"+1, "fantastic!":+5, "terrible":-2,"shit":-3, "horrible":-5}
 - [AFINN](#), [Depechemood](#), [McDonald finance](#) sentiment lexicons..
- Medium: Semantic abstractions/"clusters" per word - lemmatization, Brown clusters..
- Advanced: word2vec embeddings per word, wordnet clusters, synsets..
- Data augmentation (more data points) - translate back/forth, word2vec synonym replacement..
- Look up features about words/entities from knowledge graphs/ontologies - e.g. Wikidata/Wikipedia, Wordnet, Google




Time Series

TIME SERIES ANALYSIS

Time Series Values of Passengers



Time Series - Data over Time




Can be seen as:


- (X) Data with time
 - E.g. “Given each customers internet history, predict if they will click an Ad”
- (Y) Target over time
 - E.g.: “Evaluate a financial trading strategy, predict daily stock market price, per stock, for each day, for the next month”
 - How many coronavirus patients..

Time-series: Some vital statistics

- Stationary (over time)?
 - Do statistical properties (e.g. mean, VAR) change? (e.g. inflation, growth)
 - Dickey–Fuller test (ADF)
 - If non stationary over time - detrend!
<https://people.duke.edu/~rnau/411diff.htm>
<http://people.duke.edu/~rnau/whatuse.htm>

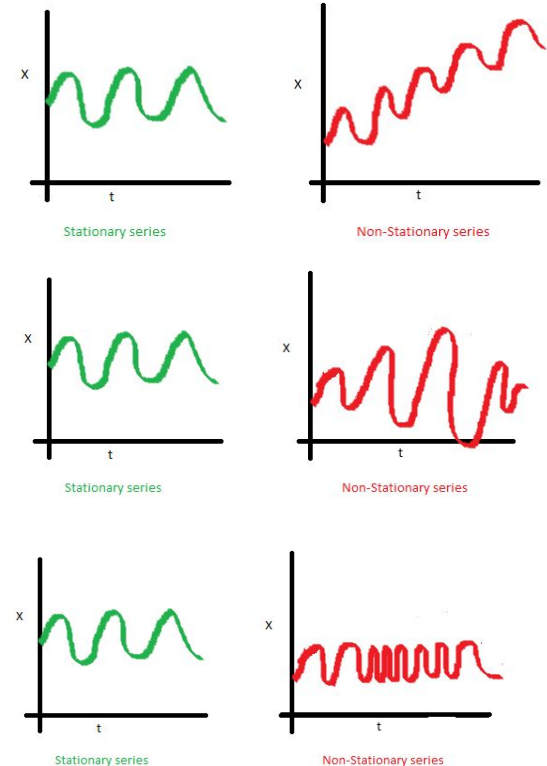


Time differencing




What does being stationary mean?

1. The mean of the series should not be a function of time. The red graph below is not stationary because the mean increases
2. The **variance** of the series should not be a function of time. This is known as **homoscedasticity**. Notice the varying **spread** of data over time.
3. Finally, the **covariance** of the i th term and the $(i + m)$ th term should not be a function of time. In the graph, notice the spread becomes closer as the time increases.



Time-series - vital stats

- Seasonality?
 - Seasonal components: Daily/monthly/hourly? Holiday?
- Noise?
 - Regularly sampled intervals? Future non-causal noise?
- Multiple variables or univariate?
 - Predict a stock by its history: univariate (+AutoRegression)
 - Using other stocks: multivariate
- How much history?
 - Less than a year?
- Outliers?
- Breakpoints, Level shifts?



Evaluating Forecasts: Common Gotchas

- Target/Feature leakage:
 - Calculating features without accounting for **prediction horizon**.
 - E.g. Aggregate mean of value, including future data points.
- Not splitting test-set by time
 - E.g. Stock prices.
- Forgetting simple baselines
 - Last value, mean of value..
 - Overcomplicated models (often inferior to naive baselines!)
- Assuming everything is predictable
- Not accounting for non stationary values (Trends..)
 - Make it stationary by 1st/2d order differencing! <https://people.duke.edu/~rnau/411diff.htm>
 - Random forests & regression: can't predict target outside of range

Simple Time-Series Features

DateTime/Calendar features

```
df[“Datetime”] = pd.to_datetime(df[“Datetime”],  
infer_datetime_format=True)
```

```
df[“interval”] = df[‘end_date’]-df[‘start_date’]
```

```
df[‘year’] = df[‘Datetime’].dt.year
```

```
df[‘month’] = df[‘Datetime’].dt.month
```

```
df[‘week’] = df[‘Datetime’].dt.week
```

```
df[‘day’] = df[‘Datetime’].dt.day
```

```
df[‘hour’] = df[‘Datetime’].dt.hour
```

```
df[‘dayofweek’] = df[‘Datetime’].dt.dayofweek
```

Datetime Properties

<code>Series.dt.date</code>	Returns numpy array of python datetime date objects (namely, the date part of Timestamps without timezone information).
<code>Series.dt.time</code>	Returns numpy array of datetime time.
<code>Series.dt.year</code>	The year of the datetime
<code>Series.dt.month</code>	The month as January=1, December=12
<code>Series.dt.day</code>	The days of the datetime
<code>Series.dt.hour</code>	The hours of the datetime
<code>Series.dt.minute</code>	The minutes of the datetime
<code>Series.dt.second</code>	The seconds of the datetime
<code>Series.dt.microsecond</code>	The microseconds of the datetime
<code>Series.dt.nanosecond</code>	The nanoseconds of the datetime
<code>Series.dt.week</code>	The week ordinal of the year
<code>Series.dt.weekofyear</code>	The week ordinal of the year
<code>Series.dt.dayofweek</code>	The day of the week with Monday=0, Sunday=6
<code>Series.dt.weekday</code>	The day of the week with Monday=0, Sunday=6
<code>Series.dt.dayofyear</code>	The ordinal day of the year
<code>Series.dt.quarter</code>	The quarter of the date
<code>Series.dt.is_month_start</code>	Logical indicating if first day of month (defined by frequency)
<code>Series.dt.is_month_end</code>	Indicator for whether the date is the last day of the month
<code>Series.dt.is_quarter_start</code>	Indicator for whether the date is the first day of a quarter.
<code>Series.dt.is_quarter_end</code>	Indicator for whether the date is the last day of a quarter.
<code>Series.dt.is_year_start</code>	Indicate whether the date is the first day of a year.
<code>Series.dt.is_year_end</code>	Indicate whether the date is the last day of the year.
<code>Series.dt.is_leap_year</code>	Boolean indicator if the date belongs to a leap year.
<code>Series.dt.daysinmonth</code>	The number of days in the month
<code>Series.dt.days_in_month</code>	The number of days in the month
<code>Series.dt.tz</code>	
<code>Series.dt.freq</code>	

Source:

https://pandas.pydata.org/docs/user_guide/timeseries.html#time-date-components

<https://www.kaggle.com/danofer/datetime-embeddings-for-end-to-end-deep-learning>

Lag

- Variable's value, X “steps” ago.
- Strong baseline to beat.
 - Momentum strategy in stocks: “Stock will be the same as yesterday” (Lag 1 day)
 - Weather will be the same as it was last year (lag 365)


```
df[“value_lag1”] = df[“Value”].shift(1)
```

Date	Value	Value _{t-1}	Value _{t-2}
1/1/2017	200	NA	NA
1/2/2017	220	200	NA
1/3/2017	215	220	200
1/4/2017	230	215	220
1/5/2017	235	230	215
1/6/2017	225	235	230
1/7/2017	220	225	235
1/8/2017	225	220	225
1/9/2017	240	225	220
1/10/2017	245	240	225

Lag: How to pick lags?

1. Domain knowledge:
 - Store sales: same weekday last week (lag7) ; last month (lag30), last year (lag 365)
2. Partial autocorrelation plot: pick points with highest (absolute) correlation with target:
 - `Pandas.plotting.autocorrelation_plot`

```
from matplotlib import pyplot
from statsmodels.graphics.tsaplots import plot_pacf
series = pd.read_csv('daily-temperature.csv')
plot_pacf(series, lags=50)
pyplot.show()
```



Sliding Window Statistics

“{Statistic} Over the last X points”

- Examples:
 - **Mean** sales over last **month**
 - **Max** sales over last **year**
 - Count unique visitors in past day
 - **Sum**, Var, STD, skew, curtosis, etc' ...
- Can be combined with different window/weighting methods :


<https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.rolling.html#pandas.DataFrame.rolling>

Example: Feature of Mean sales over past 30 days, with a prediction horizon of 7 days into the future, using daily data:

```
df["mean_30day_sales"] = df["sales"].shift(7).rolling(window="30D",on="Datetime").mean()
```

Sliding Window: EWMA (exponential weighted moving average)

- Like sliding window, but give different weights to more recent points.
- E.g. average over last 3 years, but points 3 years ago have $\frac{1}{4}$ weight, 2 years ago $\frac{1}{2}$ weight..
- <https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.ewm.html>



Many more time series features...

- Variable transformations - detrend, `diff`, scale for inflation, percent change...
- Time between X1 and X2. e.g. “Time between ``date_started`` and ``current_date``”)

```
df[“interval_days_elapsed”] = (df[‘end_date’]-df[‘start_date’]).dt.days
```

- Total time elapsed (Can help “learn” inflation over time)
- Time since last occurrence of X, counts of X in last interval
- Recent X vs historical X

```
df[“monthly_sales_vs_history”] = df[“sales”].rolling(window=“30D”).mean() /  
df[“sales”].expanding().mean()
```

- X vs seasonal X - “Sales on this sunday vs previous sundays”
- Count peaks (“max”), troughs (“min”), time between min/max (local or global peak or trough)

Many more time series features...

- Signal analysis/Decomposition - FFT (Fourier transform), DWT (Discrete wavelets), Haar Wavelets...
- Autocorrelation, cross correlation with self or other time series
 - Entropy, change in autocorrelation or entropy over time...
- Many domain specific features for continuous series over time - finance
 - “Quant” features. E.g. “Candlesticks”.
 - There are many libraries with hundreds of such features. Beware of overfit..
- *Be careful of temporal leakage!!!*

[TSFresh](#) - nice time series features library

Advanced Time-Series Features


Grouped time-series

- Features for entities/subsets within groups
- Features for entities across groups/time-series
- Store sales example:
 - Sales **per item within** the store
 - Sales of an item **across** stores

```
df["mean_item_sales"] =  
df.set_index("datetime").groupby(["store","item"])["sales"].rolling(window=30).mean()
```

Classical statistical time-series models

- Can outperform ML approaches, especially on small data, noisy problems, studied domains, highly seasonal problems, etc'
- We can use these models, or combine them with ML models!
- E.g. extract Trend, seasonal components from ARIMA/ statsmodels.seasonal decomposition/ FB prophet, and use as features!



Interpolation

- Fill in missing values - required for most classical models
- This can cause a lot of bias, especially when the data is very uneven.
- I suggest only doing this when your data is from regular intervals overall, and with high temporal resolution (e.g. weekly, daily).
- Typical approaches: backfill, forward fill , fill in by average...
 - All supported in pandas

Transform temporal variables/inputs

- Differentiate/gradient, normalize by Z-score, percent change, etc'
 - E.g. Δ /rate of change in coronavirus patients (1st order diff) is more important than absolute count, in order to predict if rate (Δ) of infection is slowing or increasing
- Transform continuous variables/target to a normal distribution
 - Hit it with a log!
 - Important for many linear models!
 - [Box-cox, other power transformations](#)



Target Transformations: Change Y

- Make target stationary for modelling - e.g. subtract/divide/[differentiate](#)
 - Very important for most models!
 - Reading material in appendix
 - 1st/2d order differencing is common approach to detrending
- “Remove baseline”
 - “Subtract” mean
 - “Subtract”/divide by top feature (e.g. moving average)
 - Normalize (Z-score) by group
 - Decomposition components/forecast

<https://otexts.com/fpp2/decomposition.html>

Categorical variables & Text over time

- “Contains”.
- “Contains within last X”
- Frequency statistics (within time-window)
- Trends (e.g. breakout topics/keywords - a la Twitter)
- Word2Vec embeddings + window over time (No NN needed!)
- Order may be less important for these features, vs recency... problem dependent!
 - Workaround - looks at it as a set (ignore most ordering)
- Lots more!

Deep learning - time series

- LSTM, RNNs, Convolutional neural networks (CNN), BiLSTM+Dilated CNN + Attention, Etc'...
- In many forecasting competitions DL **lost** to classical TS approaches! But **combination** of both often **won**!



Geospatial (Geography) Features

Geospatial features - “samples near me”

Think about [predicting taxi trip travel time](#), or [housing prices](#)..

- Features about other samples “near me”:
E.g. “average price of houses nearby”, “Average price of properties within same Zipcode”...
- Consider different “sizes” to draw a shape around - K nearest, weighting by distance, etc’.
- These features can get very computationally expensive VERY easily with naive approaches!
- (My) Example notebooks with code/features:
<https://www.kaggle.com/danofer/fare-prediction-baseline-feat-eng>

Geospatial features

- Combine **Latitude** and **Longitude** columns together => new categorical feature
- Reduce Lat/Long granularity (e.g. 3d decimal place is $\sim \pm 110$ meters).

<https://gis.stackexchange.com/questions/8650/measuring-accuracy-of-latitude-and-longitude>


```
df["rounded_ll"] = df["lat"].round(3).astype(str) +  
df["lon"].round(3).astype(str)
```

- Bin/discretize Lat, Long (as categorical feature) - mainly useful for [linear or deep models](#)

Geospatial distance

Different ways of calculating distance:

- **Euclidean distance** (“as the crow flies”)
- Haversine (includes earth’s curvature), Manhattan, etc’
- Travel distance/time by car (not) - e.g. Manhattan distance, real travel time (google maps, Bing APIs), street graph distance (using OSM)

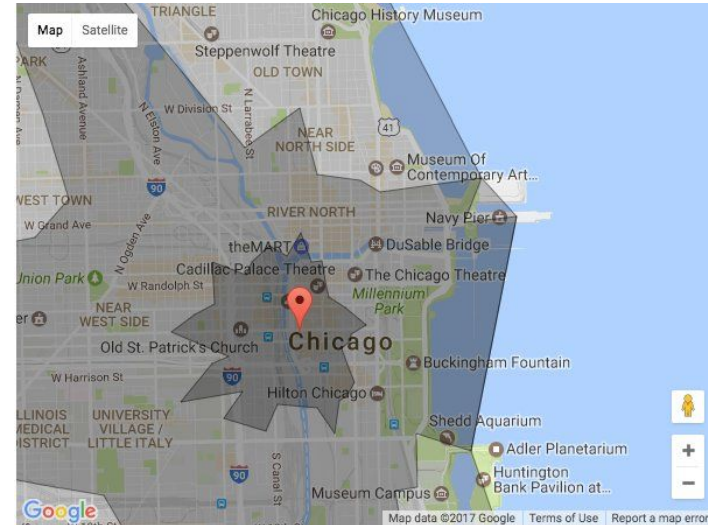


```
from math import cos, sqrt
def quick_euclid_distance(Lat1, Long1, Lat2, Long2):
    x = Lat2 - Lat1
    y = Long2 - Long1
    return sqrt(x* + y*y)
```

<https://www.kaggle.com/danofer/fare-prediction-baseline-feat-eng>

Isochrones

- **Isochrones** = “What’s within 5/15/K minutes walk/drive from me?”
 - Can be very different from “air distance” - What if there’s a highway crossing the street? Think Tel Aviv old Central bus station...
 - Use: [Open Street Map - OSMnx](#) - geoffboeing.com/2017/08/isochrone-maps-osmnx-python/ or [google maps](#)



Geospatial clusters


- Cluster lat-long points (k-means) - add as new feature or high level aggregator. Can help discover neighbourhoods.
- Use real world “clusters”! E.g. Zip codes, cities, Airbnb/Zillow business neighbourhoods, Nielsen...

Geospatial K-means Clustering by latitude and longitude

```
from sklearn.cluster import KMeans

## df = data of latitude and longitude for each row
X = df[['latitude', 'longitude']]
kmeans = KMeans(n_clusters=id_n).fit(X)
id_label=kmeans.labels_

## plot output ...
```



<https://www.kaggle.com/xxing9703/kmean-clustering-of-latitude-and-longitude>

More Geographic features

- Distance from: City center, ocean, beach, park, road, nearest other data point...
- # Points of interest nearby (OSM - [Open Street Map POI](#), [OSMnx](#))
- # tripadvisor/Yelp checkins in surrounding area, average review scores
- Geotemporal features (Great for fraud, [crime!](#)) - travel speed between points, events nearby in past hour...
- [Geocode](#) addresses into LatLong or reverse; LatLong into locations
- External data - census demographic data, ACS, Zillow, Open Street map (OSM) map, APIs



Thanks!

P.S - We're Hiring!

ddofer.github.io

Further reading

Further reading

- https://github.com/solegalli/packt_featureengineering_cookbook
- [Python Feature Engineering Cookbook](#)
- Feature Engineering for Machine Learning, O'Reilly. <https://github.com/alicezheng/feature-engineering-book>
- <https://github.com/aikho/awesome-feature-engineering>
- <https://github.com/Yimeng-Zhang/feature-engineering-and-feature-selection>
- <https://machinelearningmastery.com/basic-feature-engineering-time-series-data-python/>
- Kaggle: Learn Feature extraction/"engineering" <https://www.kaggle.com/learn/feature-engineering>
- <https://machinelearningmastery.com/basic-feature-engineering-time-series-data-python/>
- <https://www.slideshare.net/HJvanVeen/feature-engineering-72376750/11> - categorical embedding/transformations
- <https://machinelearningmastery.com/discover-feature-engineering-how-to-engineer-features-and-how-to-get-good-at-it/>
- <https://www.slideshare.net/HJvanVeen/feature-engineering-72376750/11>
- https://blog.dataiku.com/2015/06/25/predicting_crime_sf (Geospatial features)
- <https://www.kaggle.com/c/ieee-fraud-detection/discussion/108575#624914>
- Feature Engineering for Machine Learning, [Udemy Course](#)

Further reading: Time Series

- <https://people.duke.edu/~rnau/411diff.htm>
- Forecasting: Principles and Practice - Rob J Hyndman and G Athanasopoulos - <https://otexts.com/fpp2/>
- <http://www.svds.com/avoiding-common-mistakes-with-time-series/>
- [TimeSeriesSplit](#) - scikit-learn Time Series cross-validator

Further reading: Libraries

- Pandas - designed for tabular data and time-series! <https://pandas.pydata.org/>
- Automatic extraction of relevant features from time series: <http://tsfresh.readthedocs.io>
- Facebook Prophet - very easy to use bayesian TS
- Irregularly spaced TS features : <https://traces.readthedocs.io/en/latest/>
- <https://github.com/microsoft/forecasting>
- **Statsmodels** (Inc. classical models for time series - Arima)
- **OSMNx - Open street map**
- **Hyndman (R) Blog - auto.arima, Feasts etc'** - <https://robjhyndman.com/hyndsight>
- Feature Tools - Open-source Feature engineering library. Time-aware. <https://www.featuretools.com/>
- H2O - <http://docs.h2o.ai/driverless-ai/latest-stable/docs/userguide/time-series.html>
- **SparkBeyond** (Automated feature engineering & insights including complex multivariate, non numeric irregular time-series)